

COURSEWORK 2

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

Computational Linear Algebra

Author:

Richard Fu (CID: 01852979)

Date: December 8, 2022

1 QUESTION 1

1.1 Part a

The function that generates the matrix $A^{(n)}$ as described is located in `cw2/q1.py`, and is called `generate_A()`. In the same file, there is a function called `get_rho()` which can be called on a matrix to find the growth factor ρ as defined in section 4.3. Below these functions is the code that creates $A^{(6)}$, and prints the corresponding growth factor ρ , which has a numerical value of 32.0.

1.2 Part b

The LU factorisation of $A^{(n)}$ is given by $A^{(n)} = LU$, where:

$$L = \begin{bmatrix} 1 & 0 & \dots & 0 & 0 \\ -1 & 1 & \ddots & \vdots & \vdots \\ -1 & -1 & \ddots & 0 & \vdots \\ \vdots & \vdots & \ddots & 1 & 0 \\ -1 & -1 & \dots & -1 & 1 \end{bmatrix}, \quad U = \begin{bmatrix} 1 & 0 & \dots & 0 & 2^0 \\ 0 & 1 & \ddots & \vdots & 2^1 \\ 0 & 0 & \ddots & 0 & \vdots \\ \vdots & \vdots & \ddots & 1 & 2^{n-2} \\ 0 & 0 & \dots & 0 & 2^{n-1} \end{bmatrix}. \quad (1)$$

To see this, first notice that the first $(n-1)$ columns of U are $(e_1, e_2, \dots, e_{n-1})$, where e_i is the i -th canonical basis vector. This means that the first $(n-1)$ columns of LU are just the first $(n-1)$ columns of L .

Now, let us look at the entries in the n -th column of LU .

$$(LU)_{in} = \begin{cases} 2^0 & i = 1 \\ 2^{i-1} - \sum_{k=0}^{i-2} 2^k & i \geq 2. \end{cases} \quad (2)$$

Trivially $2^0 = 1$, and also $2^{i-1} - \sum_{k=0}^{i-2} 2^k = 2^{i-1} - \frac{1-2^{i-1}}{1-2} = 2^{i-1} + 1 - 2^{i-1} = 1$.

Therefore, LU is given by L with the last column changed to a vector of ones, which is precisely $A^{(n)}$.

$\rho = \frac{\max_{ij} |u_{ij}|}{\max_{ij} |a_{ij}|} = \frac{2^{n-1}}{1} = 2^{n-1}$, so the growth factor for $A^{(n)}$ does depend on n .

1.3 Part c

The functions `error_LUP()` and `error_solve_LUP()` located in `cw2/q1.py` help to calculate the errors from the LUP factorisation, and the errors in the least squares solution to $Ax = b$ respectively. Underneath we call the functions on $A^{(60)}$, and we see that the errors are quite large for both.

1.4 Part d

The function `random_matrix()` in `cw2/q1.py` helps to generate random $n \times n$ matrices with entries sampled from a $\text{Uniform}(-\frac{1}{n}, \frac{1}{n})$ distribution.

Located in `cw2/plotting_scripts/q1_plots.py` is the code that calculates the growth factor ρ of a random matrix created by `random_matrix()`, for a range of n values, and then generates the plot with log-scales we can see in Figure 1. We can see that the black scatter points seem to follow a straight line quite well, which suggests that there is a power relationship between ρ and n ; i.e. $\rho = an^k$. In fact, the red line we see is actually the curve $\rho = \frac{1}{3}n^{\frac{3}{4}}$, and this appears to be a good fit for the scatter points for the random matrices generated.




Figure 1: Scatter plot of growth values ρ for varying of n

This investigation shows that in general we will not get an exponential relationship between matrix dimension n and growth factor ρ like we did for matrices $A^{(n)}$, so it is not all bad news for using Gaussian elimination.

2 QUESTION 2

2.1 Part a

The function `MGS_solve_ls()` is located in `cw2/q2.py`. The pytests corresponding to this function are located in `cw2/test2.py` and are called by running `test_MGS_solve_ls()`.

2.2 Part b

In `cw2/plotting_scripts_q2_plots.py` there is code that repeatedly generates random x^* and calculates the numerical error when solving the least squares problem using `MGS_solve_ls()`, and also when using `householder_ls()` from `cla_utils`. We then plot these errors with log y-axis scale as seen in Figure 2.




Figure 2

Clearly, we can see that the `householder_ls()` function outperforms `MGS_solve_ls()` with respect to error in numerical least squares solution.

2.3 Part c

First,

$$A_+ = \begin{bmatrix} A & | & b \end{bmatrix}. \quad (3)$$

$$A_+ = \begin{bmatrix} \hat{Q} & | & q_{n+1} \end{bmatrix} \begin{bmatrix} \hat{R} & | & z \\ 0 & | & \rho \end{bmatrix}, \quad (4)$$

$$= \begin{bmatrix} \hat{Q}\hat{R} & | & \hat{Q}z + \rho q_{n+1} \end{bmatrix}. \quad (5)$$

$$(3) + (5) \implies b = \hat{Q}z + \rho q_{n+1} \quad (6)$$

Now to find an equation for z , we left multiply (4) by $[\hat{Q} \mid \mathbf{q}_{n+1}]^*$:

$$[\hat{Q} \mid \mathbf{q}_{n+1}]^* [A \mid \mathbf{b}] = \begin{bmatrix} \hat{R} & z \\ \mathbf{0} & \rho \end{bmatrix}, \quad (7)$$

$$\Rightarrow \begin{bmatrix} \hat{Q}^* \\ \mathbf{q}_{n+1}^* \end{bmatrix} [A \mid \mathbf{b}] = \begin{bmatrix} \hat{R} & z \\ \mathbf{0} & \rho \end{bmatrix}, \quad (8)$$

$$\Rightarrow \begin{bmatrix} \hat{Q}^* A & \hat{Q}^* \mathbf{b} \\ \mathbf{q}_{n+1}^* A & \mathbf{q}_{n+1}^* \mathbf{b} \end{bmatrix} = \begin{bmatrix} \hat{R} & z \\ \mathbf{0} & \rho \end{bmatrix}, \quad (9)$$

$$\Rightarrow z = \hat{Q}^* \mathbf{b}. \quad (10)$$

2.4 Part d

Our modified algorithm proceeds as follows:

1. First form $A_+ = [A \mid \mathbf{b}]$.
2. Find the reduced QR factorisation $[\hat{Q} \mid \mathbf{q}_{n+1}] \begin{bmatrix} \hat{R} & z \\ \mathbf{0} & \rho \end{bmatrix} = A_+$ using our modified Gram-Schmidt implementation `GS_modified()` from `cla_utils`.
3. Slice out both the vector $z = \hat{Q}^* \mathbf{b}$, and the matrix \hat{R} .
4. Solve $\hat{R}\mathbf{x} = z$ using our backward substitution implementation `solve_U()` from `cla_utils`.

The operation count for modified Gram-Schmidt is $\mathcal{O}(mn^2)$, and for backward substitution it is $\mathcal{O}(n^2)$. Therefore, the operation count for our modified algorithm is dominated by modified Gram-Schmidt, and the operation count is $\mathcal{O}(mn^2)$.

The function `MGS_solve_ls_modified()` is located in `cw2/q2.py`, and the corresponding pytests are called by running `test_MGS_solve_ls_modified()` in `cw2/test2.py`.

2.5 Part e

The code applying `MGS_solve_ls_modified()` to the provided matrix A_2 is also located in `cw2/plotting_scripts/q2_plots.py`. We then plot the errors in the numerical least squares solution using all three algorithms, which can be seen in Figure 3. Again, we have used a log y-axis scale.




Figure 3

From Figure 3 we can clearly see that `MGS_solve_ls_modified()` performs very similarly to the `householder_ls()` and is a major improvement from `MGS_solve_ls()`, with respect to error in numerical least squares solution.

2.6 Part f

The code that generates \mathbf{b} outside the range space of A_2 , and calculates the errors in numerical least squares solution using this \mathbf{b} is located in `cw2/plotting_scripts/q2_plots.py`. We then plot the errors in the numerical least squares solution using this \mathbf{b} using all three algorithms, which can be seen in Figure 4. Again, we have used a log y-axis scale.

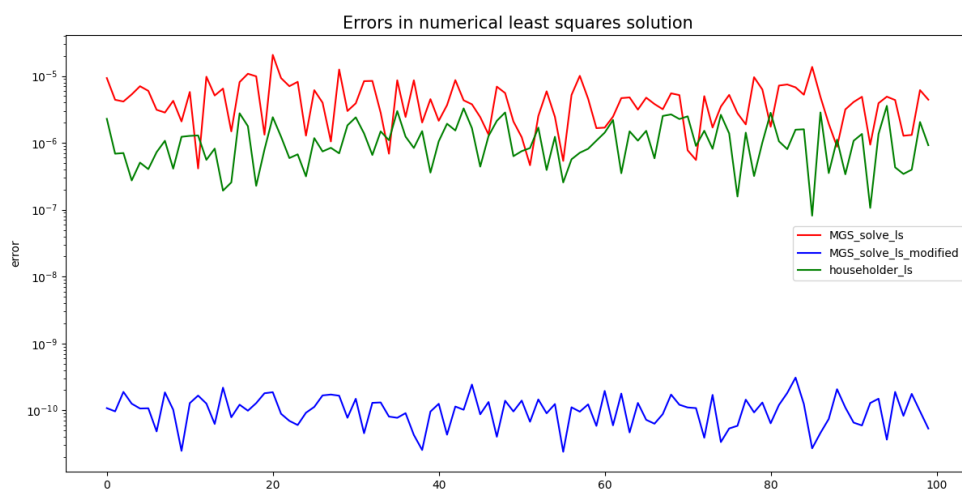


Figure 4

From Figure 4, it appears that the improvement in `MGS_solve_ls_modified()` is evident even in this case of \mathbf{b} outside the range space of A_2 , as it performs better than both `MGS_solve_ls()` and `householder_ls()` with respect to error in numerical least squares solution.

3 QUESTION 3

3.1 Part a

The formula for the $N^2 \times N^2$ matrix D is given below:

$$D_{ij} = \begin{cases} 1 + 4s^2 & i = j \\ -s^2 & |j - i| = N \\ -s^2 & j - i = 1, \text{ and } i \not\equiv 0 \pmod{N} \\ -s^2 & i - j = 1, \text{ and } j \not\equiv 0 \pmod{N} \\ 0 & \text{otherwise.} \end{cases} \quad (11)$$

The bandwidth of matrix D is N .

3.2 Part b

The function `LU_inplace()` in `cla_utils/exercises6.py` has been updated to support the banded version of LU factorisation. Lower bandwidth can be specified by keyword argument `bl`, and upper bandwidth can be specified by keyword argument `bu`. The corresponding pytest tests that test this implementation are located in `cw2/test3.py`, and can be called by running `test_LU_inplace_banded()`.

The function `solve_L()` has been updated to support the more efficient forward substitution algorithm for banded lower triangular matrices, and this lower bandwidth can be specified by keyword argument `bl`. The corresponding pytest tests that test this implementation are located in `cw2/test3.py`, and can be called by running `test_solve_L_banded()`.

The function `solve_U()` has been updated to support the more efficient backward substitution algorithm for banded upper triangular matrices, and this upper bandwidth can be specified by keyword argument `bu`. The corresponding pytest tests that test this implementation are located in `cw2/test3.py`, and can be called by running `test_solve_U_banded()`.

Our function `solve_LU()` is located in `cw2/q3.py`, and can take keyword arguments `bl` (lower bandwidth), and `bu` (upper bandwidth). If we are solving a banded matrix system, we can specify the bandwidths to utilise the bandwidth support added to the functions. The corresponding pytest tests that test this implementation are located in `cw2/test3.py`, and can be called by running `test_solve_LU_banded()`.

3.3 Part c

The function `image_simulator()` is located in `cw2/q3.py`, and there is also a helper function called `construct_D()` that returns D as defined above, and this is called in the `image_simulator()` function. There is also a keyword argument `banded_solve`

which specifies whether to use the banded solver.

In `cw2/plotting_scripts/q3_plots.py` is the code that draws the plots in Figure 5.

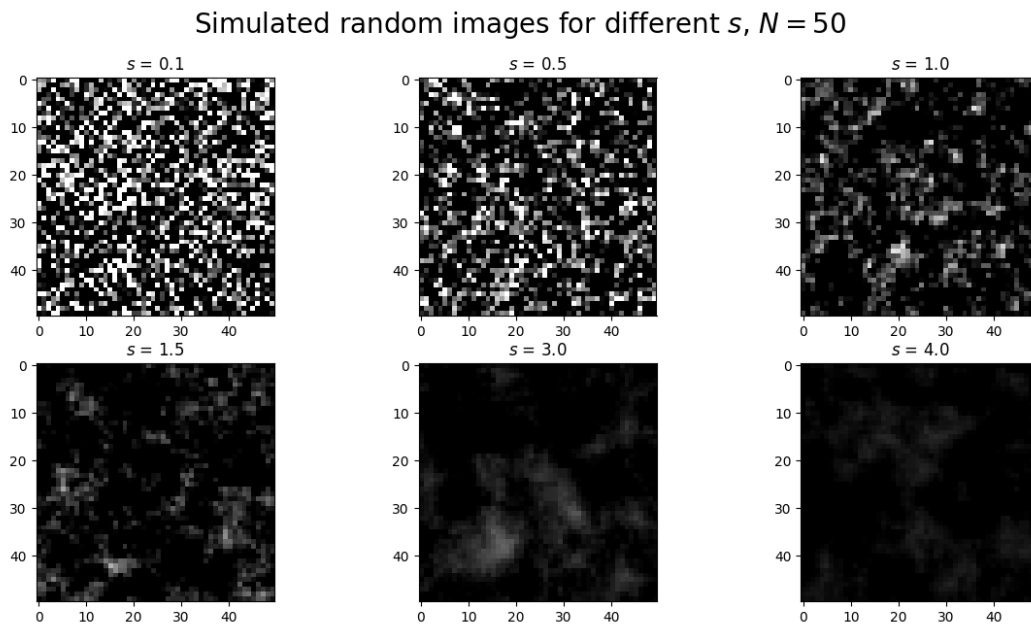


Figure 5

In Figure 5, we can see that the random simulated images become darker as s increases. s is a parameter that governs the strength of correlations between nearby points, so this behaviour is as we expect since we assume the points outside of the grid are black.

3.4 Part d

The function `image_simulator()` in `cw2/q3.py` has a `return_time` keyword argument that specifies whether to return the runtime of the simulation rather than the actual array. We utilise this in the code in `cw2/plotting_scripts/q3_plots.py` to obtain our runtimes for different N , and draw the plot in Figure 6.

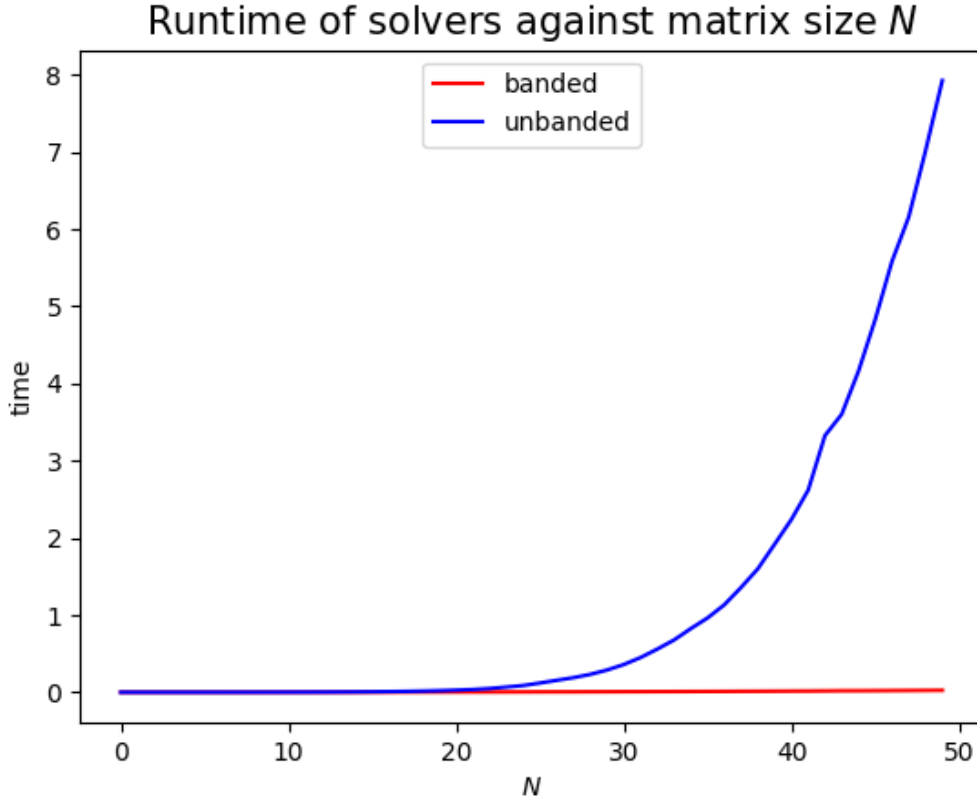


Figure 6

From Figure 6, we can see that the runtime for the unbanded solver seems to grow exponentially with N , whilst the runtime for the banded solver is very fast for all values of N . \mathbf{D} is size N^2 , and has lower and upper bandwidth size equal to N . This means that the operation count for our banded solver is $\mathcal{O}(N^4)$, whilst the operation count for our unbanded solver is $\mathcal{O}(N^6)$. This is consistent with the observed timings.

3.5 Part e

Our iterative equations are:

$$((1 + \rho)\mathbf{I} + s^2\mathbf{H})\hat{\mathbf{u}}^{n+\frac{1}{2}} = (\rho\mathbf{I} - s^2\mathbf{V})\hat{\mathbf{u}}^n + \mathbf{w}, \quad (12)$$

$$((1 + \nu)\mathbf{I} + s^2\mathbf{V})\hat{\mathbf{u}}^{n+1} = (\nu\mathbf{I} - s^2\mathbf{H})\hat{\mathbf{u}}^{n+\frac{1}{2}} + \mathbf{w}. \quad (13)$$

Let us assume the algorithm converges to a limit \mathbf{u}^* , i.e. $\hat{\mathbf{u}}^{n+1} = \hat{\mathbf{u}}^{n+\frac{1}{2}} = \hat{\mathbf{u}}^n = \mathbf{u}^*$. Our iterative equations now become:

$$((1 + \rho)\mathbf{I} + s^2\mathbf{H})\mathbf{u}^* = (\rho\mathbf{I} - s^2\mathbf{V})\mathbf{u}^* + \mathbf{w}, \quad (14)$$

$$((1 + \nu)\mathbf{I} + s^2\mathbf{V})\mathbf{u}^* = (\nu\mathbf{I} - s^2\mathbf{H})\mathbf{u}^* + \mathbf{w}. \quad (15)$$

$$(14) \implies (\mathbf{I} + s^2\mathbf{V} + s^2\mathbf{H})\mathbf{u}^* = \mathbf{w}. \quad (16)$$

$$(15) \implies (\mathbf{I} + s^2 \mathbf{V} + s^2 \mathbf{H}) \mathbf{u}^* = \mathbf{w}. \quad (17)$$

Solving this equation:

$$(P(\mathbf{I} + s^2 \mathbf{V} + s^2 \mathbf{H}) \mathbf{u}^*)_{i,j} = (P \mathbf{u}^*)_{i,j} + s^2 (P \mathbf{V} \mathbf{u}^*)_{i,j} + s^2 (P \mathbf{H} \mathbf{u}^*)_{i,j}, \quad (18)$$

$$= u_{i,j}^* + s^2 (2u_{i,j}^* - u_{i,j-1}^* - u_{i,j+1}^*) + s^2 (2u_{i,j}^* - u_{i-1,j}^* - u_{i+1,j}^*), \quad (19)$$

$$= u_{i,j}^* + s^2 (-u_{i-1,j}^* - u_{i+1,j}^* - u_{i,j-1}^* - u_{i,j+1}^* + 4u_{i,j}^*), \quad (20)$$

$$= (P \mathbf{D} \mathbf{u}^*)_{i,j}, \quad (21)$$

$$\implies (\mathbf{I} + s^2 \mathbf{V} + s^2 \mathbf{H}) = \mathbf{D}, \quad (22)$$

$$\implies \mathbf{D} \mathbf{u}^* = \mathbf{w}. \quad (23)$$

3.6 Part f

First, it is important to note that \mathbf{H} is an $N^2 \times N^2$ banded matrix with 2's on its main diagonal, and -1's on the N-th super and sub diagonals.

Let us introduce the notation $\hat{\mathbf{x}}_{perm}$ to denote the order 2 permutation of an N^2 -length vector $\hat{\mathbf{x}}$, such that $(P\hat{\mathbf{x}}) = (P\hat{\mathbf{x}}_{perm})^T$, and $(P\hat{\mathbf{x}}_{perm}) = (P\hat{\mathbf{x}})^T$, where P is the mapping defined in the question.

Now note that:

$$\mathbf{H} \hat{\mathbf{x}} = (\mathbf{V} \hat{\mathbf{x}}_{perm})_{perm} \iff (\mathbf{H} \hat{\mathbf{x}})_{perm} = \mathbf{V} \hat{\mathbf{x}}_{perm}, \quad (24)$$

$$\implies \mathbf{H} \hat{\mathbf{x}} = \mathbf{b} \iff \mathbf{V} \hat{\mathbf{x}}_{perm} = \mathbf{b}_{perm}. \quad (25)$$

It then follows:

$$((1 + \rho)\mathbf{I} + s^2 \mathbf{H}) \hat{\mathbf{u}}^{n+\frac{1}{2}} = (\rho \mathbf{I} - s^2 \mathbf{V}) \hat{\mathbf{u}}^n + \mathbf{w}, \quad (26)$$

$$\iff \underbrace{((1 + \rho)\mathbf{I} + s^2 \mathbf{V})}_{\mathbf{A}_1} \underbrace{\hat{\mathbf{u}}_{perm}^{n+\frac{1}{2}}}_{\mathbf{x}_1} = \underbrace{((\rho \mathbf{I} - s^2 \mathbf{V}) \hat{\mathbf{u}}^n + \mathbf{w})}_{\mathbf{b}_1} \Big|_{perm}. \quad (27)$$

$$((1 + \nu)\mathbf{I} + s^2 \mathbf{V}) \hat{\mathbf{u}}^{n+1} = (\nu \mathbf{I} - s^2 \mathbf{H}) \hat{\mathbf{u}}^{n+\frac{1}{2}} + \mathbf{w}, \quad (28)$$

$$\iff \underbrace{((1 + \nu)\mathbf{I} + s^2 \mathbf{V})}_{\mathbf{A}_2} \underbrace{\hat{\mathbf{u}}^{n+1}}_{\mathbf{x}_2} = \underbrace{((\nu \mathbf{I} - s^2 \mathbf{V}) \hat{\mathbf{u}}_{perm}^{n+\frac{1}{2}})}_{\mathbf{b}_2} \Big|_{perm} + \mathbf{w}. \quad (29)$$

Our systems (27) and (29) both involve banded matrices \mathbf{A}_1 and \mathbf{A}_2 with bandwidths 1, which can be used to our advantage in terms of efficiency.

Also, \mathbf{A}_1 is a block diagonal matrix with N repeated identical blocks of size N that also must have bandwidth 1, which we will denote \mathbf{A}'_1 .

We can view our system involving A_1 and \mathbf{b}_1 as equivalent to solving $A_1' \mathbf{x}_1^{(i)} = \mathbf{b}_1^{(i)}$, for $i = 1, \dots, N$, where $\mathbf{b}_1 = \begin{bmatrix} \mathbf{b}_1^{(1)} \\ \vdots \\ \mathbf{b}_1^{(N)} \end{bmatrix}$, $\mathbf{x}_1 = \begin{bmatrix} \mathbf{x}_1^{(1)} \\ \vdots \\ \mathbf{x}_1^{(N)} \end{bmatrix}$.

An identical argument applies to the system involving A_2 , \mathbf{b}_2 , and \mathbf{x}_2 .

Our algorithm for each iteration is as follows:

1. Evaluate the RHS of equation (27), \mathbf{b}_1 , and use `numpy.reshape()` to cast this to the form $\mathbf{B}_1 = \begin{bmatrix} \mathbf{b}_1^{(1)} & \dots & \mathbf{b}_1^{(N)} \end{bmatrix}$.
2. Generate A_1' , the size N diagonal block of A_1 .
3. Use our `solve_LU()` function from `cw2/q3.py` to solve $A_1' \mathbf{x}_1^{(i)} = \mathbf{b}_1^{(i)}$, for $i = 1, \dots, N$ simultaneously, with keyword arguments `bl = 1` and `bu = 1`. This returns \mathbf{x}_1 cast to a matrix.
4. Evaluate the RHS of equation (29), \mathbf{b}_2 , and again use `solve_LU()` function from `cw2/q3.py` to solve $A_2' \mathbf{x}_2^{(i)} = \mathbf{b}_2^{(i)}$, for $i = 1, \dots, N$ simultaneously, with keyword arguments `bl = 1` and `bu = 1`. This returns \mathbf{x}_2 cast to a matrix.
5. Use `numpy.flatten()` to flatten this matrix column-wise to obtain \mathbf{x}_2 .

In steps 1 and 4, when performing the multiplication of the bandwidth 1 matrix multiplied by another matrix, we can take advantage of the bandedness so that the operation count is $\mathcal{O}(N^2)$ rather than the standard $\mathcal{O}(N^3)$. The operation counts for solving the two band-1, size N systems are both $\mathcal{O}(N^2)$, so the whole algorithm is $\mathcal{O}(N^2)$.

This is less than the operation count for the direct solving of the band- N system, which is $\mathcal{O}(N^4)$.

It requires less memory, as utilising the block diagonal properties of A_1 and A_2 means that we only have to store $N \times N$ matrices rather than $N^2 \times N^2$ matrices.

3.7 Part g

All the functions mentioned are located in `cw2/q3.py`. The function `construct_V_block()` is used to help construct one of the blocks of V .

The function `mat_mul_banded()` is an implementation of the optimised matrix multiplication involving a banded matrix as described above. The corresponding pytests that make sure this is working properly are located in `cw2/test3.py`, and are called by running `test_mat_mul_banded()`.

There is also a `run_iter()` function that performs the algorithm described in 3.6, and this is called by the `iterative_solve()` function iteratively until our stopping condition $\|\mathbf{D}\hat{\mathbf{u}}^n - \mathbf{w}\| < \epsilon\|\mathbf{w}\|$ is met. The corresponding pytests are called by running `test_run_iter()` and `test_iterative_solve()`, which are also located in `cw2/test3.py`.

3.8 Part h

All the code for this question is located in `cw2/plotting_scripts/q3_plots.py`. In order to explore the effects of ρ and ν on the number of iterations we calculate the number of iterations for combinations of ρ and ν between 0.2 and 2 at 0.1 intervals. We sort these values into a matrix, and produce a heat map for various N values. The heat map should be darker for lower iteration counts, and lighter for higher iteration counts.

Heat Map plot of iterations for varying N




Figure 7

In Figure 7 we see that for all N , the optimal values regions for ρ and ν both seem to be centered around 1. This optimal value region appears to converge to $\rho = 1$, $\nu = 1$ as N increases, so it suggests we should take these our optimal values.




Figure 8

In Figure 8 we see our typical iteration count is 9 using our optimal values and $\epsilon = 1.0e - 6$, and this does not appear to depend on N . This suggests our operation count is still just $\mathcal{O}(N^2)$.

3.9 Part i

With a parallel computer that can execute multiple operations at once (multi-core), we can execute the solving of the set of systems $A_1' x_1^{(i)} = b_1^{(i)}$, for $i = 1, \dots, N$ simultaneously rather than one after the other as our current implementation does, and similarly for the second set of systems $A_2' x_2^{(i)} = b_2^{(i)}$, for $i = 1, \dots, N$. As long as the computer performance is unaffected by the simultaneous operations, this will accelerate the runtime of algorithm.